

## Model Development Phase

**Project Name :** Amazon Instrument Reviews

### Model Selection Report

In this project, multiple machine learning models were evaluated for sentiment classification of Amazon Musical Instrument Reviews. Factors such as accuracy, confusion matrix, classification report were considered to select the most suitable model.

Model	Description
Model - 1 : Logistic Regression	<ul style="list-style-type: none"><li>· A simple linear model applied to TF-IDF features.</li><li>· Served as a baseline for sentiment classification.</li></ul>
Model - 2 : Random Forest	<ul style="list-style-type: none"><li>· An ensemble learning method with multiple decision trees.</li><li>· Achieved higher accuracy and robustness compared to Logistic Regression. Selected as the final model.</li></ul>

Model	Description
-------	-------------

## Model 1 – Training Setup & Visualizations

```
import numpy as np
import pandas as pd
```

```
import zipfile
import os

# Extract both files from the zip
with zipfile.ZipFile('archive.zip', 'r') as zip_ref:
    # List all files in the zip
    print("Files in zip:")
    for file_info in zip_ref.filelist:
        print(f"- {file_info.filename}")

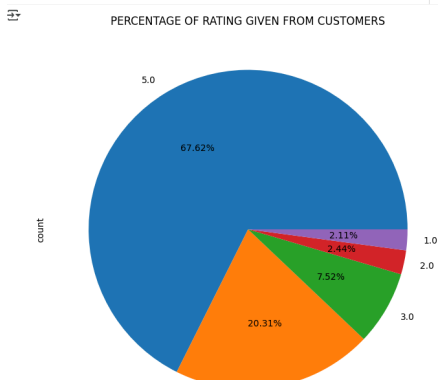
    # Extract both files
    zip_ref.extractall('extracted_files')

print("Files extracted successfully!")
```

```
Files in zip:
- Musical_Instruments_5.json
- Musical_instruments_reviews.csv
Files extracted successfully!
```

```
#percentage of rating given from customers
from collections import Counter
import matplotlib.pyplot as plt

df.overall.value_counts().plot(kind = "pie", legend = False, autopct = "%1.2f%%", fontsize = 10, figsize = (8,8))
plt.title("PERCENTAGE OF RATING GIVEN FROM CUSTOMERS", loc = "center")
plt.show()
```




#Labeling product based on Ratings

```
def Labelling(Rows) :
    if(Rows["overall"] > 3.0) :
        Label = "Positive"
    elif(Rows["overall"] < 3.0) :
        Label = "Negative"
    else :
        Label = "Neutral"

    return Label
```

```
df["sentiment"] = df.apply(Labeling, axis = 1)
df["sentiment"].value_counts().plot(kind = 'bar', color = "blue")
plt.title("Rating")
plt.xlabel("Sentiments", color = "green", fontsize = 10, labelpad = 15)
plt.xticks(rotation = 0)
plt.ylabel("Amount of sentiments")
plt.show()
```



## Model 2 – Dataset Loading & Balancing

```
# Text cleaning
import re
import string

def Text_cleaning(Text):

    Text = Text.lower() # 1: lower case
    punc = str.maketrans(' ', '', string.punctuation)
    Text = Text.translate(punc) # 2: Remove punctuation
    Text = re.sub(r'\d+', '', Text) # 3: Remove numbers
    Text = re.sub(r'https?://\S+|www\.\S+', '', Text) # 4: Remove links
    Text = re.sub('\n', '', Text) # 5: Delete new lines

    return Text
```

!pip install nltk

```
Requirement already satisfied: nltk in /usr/local/lib/python3.12/dist-packages (3.9.1)
Requirement already satisfied: Click in /usr/local/lib/python3.12/dist-packages (from nltk) (8.2.1)
Requirement already satisfied: joblib in /usr/local/lib/python3.12/dist-packages (from nltk) (1.5.2)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.12/dist-packages (from nltk) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from nltk) (4.67.1)
```

```
# Text Preprocessing
import nltk
import nltk.corpus
from nltk.corpus import stopwords # Import stopwords here
nltk.download("punkt")
nltk.download("stopwords")
nltk.download("wordnet")
nltk.download("punkt_tab")
from nltk.stem import WordNetLemmatizer
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip...
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip...
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Unzipping package punkt_tab to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt_tab.zip...
```

```
# Text Preprocessing function
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from nltk.tokenize import word_tokenize

def Text_Preprocessing(Text):
    # Tokenization
    tokens = word_tokenize(Text)

    # Remove stop words
    stop_words = set(stopwords.words('english'))
    tokens = [word for word in tokens if word not in stop_words]

    # Lemmatization
    lemmatizer = WordNetLemmatizer()
    tokens = [lemmatizer.lemmatize(word) for word in tokens]

    # Join tokens back into a string
    return " ".join(tokens)
```

```
df["review"] = df["review"].apply(lambda Text : Text_cleaning(Text))
df["review"] = df["review"].apply(lambda Text : Text_Preprocessing(Text))
df.head()
```

	reviewID	asin	reviewerName	helpful	overall	unhelpful	reviewTime	reviewText	reviews	sentiment
0	A2EP820J2R5U	1384719342	cassandra	to "Yeah, well, that's just like, u..."	[5, 5]	5.0	1302545600	02/28, 2014	much write exactly supposed filter pop sound c...	Positive
1	A1A4N7EAXKD8S	1384719342	Jake	[13, 14]	5.0	1363392000	03/16, 2013	product exactly quite affordable realized dou...	Positive	


```

from textblob import TextBlob

df["Polarity"] = df["reviews"].map(lambda Text : TextBlob(Text).polarity)
df["Polarity"].plot(kind="hist", bins=20, edgecolor="black", linewidth=1, color="orange", figsize=(10,5))

plt.title("Polarity Score in Reviews", fontsize=15, pad=20)
plt.xlabel("Polarity", labelpad=5, color="red")
plt.ylabel("Amount of Reviews", labelpad=20, color="green")
plt.show()

```




```

# Review length
df["length"] = df["reviews"].apply(lambda x: len(x.split()))

# Use plt.hist directly
plt.figure(figsize=(10,5))
plt.hist(df["length"], bins=20, edgecolor="blue", linewidth=1, color="orange")
plt.title("Length of Reviews", color="blue", pad=20)
plt.xlabel("Length", labelpad=15, color="red")
plt.ylabel("Frequency")
plt.show()

```




```

# Word Counts
df["Word_count"] = df["reviews"].apply(lambda x: len(x.split()))
plt.hist(df["Word_count"])

(array([9.486e+03, 6.410e+02, 1.500e+02, 3.900e+01, 1.000e+01, 7.000e+00,
        4.000e+00, 1.000e+00, 2.000e+00, 1.000e+00]),
array([ 2. , 110.4, 218.8, 327.2, 435.6, 544. , 652.4, 760.8,
        869.2, 977.6, 1086. ]),
<BarContainer object of 10 artists>)

```



	<pre> # N-Gram Analysis def GramAnalysis(Corpus, Gram, N):     Vectorizer = CountVectorizer(stop_words="english", ngram_range=(Gram, Gram))      ngram_matrix = Vectorizer.fit_transform(Corpus)      # N-Gram Frequency     Counts = ngram_matrix.sum(axis=0)      # List of words     words = [(word, Counts[0, idx]) for word, idx in Vectorizer.vocabulary_.items()]      # Sort Descending     words = sorted(words, key=lambda x: x[1], reverse=True)      return words[:N]  # Filter the platforms Based on Sentiments Positive = df[df["sentiment"]=="Positive"].dropna() Negative = df[df["sentiment"]=="Negative"].dropna() Neutral = df[df["sentiment"]=="Neutral"].dropna() </pre>
<h2>Model 3 – Feature Extraction (TF-IDF)</h2>	<pre> #Unigram of reviews based on sentiment from sklearn.feature_extraction.text import CountVectorizer  words = GramAnalysis(Positive["reviews"], 1, 20) Unigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Unigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Unigram of reviews with Positive Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  # Negative words = GramAnalysis(Negative["reviews"], 1, 20) Unigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Unigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Unigram of reviews with Negative Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  # Neutral words = GramAnalysis(Neutral["reviews"], 1, 20) Unigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Unigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Unigram of reviews with Neutral Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  # Bigram - Positive, Negative, Neutral  #Positive words = GramAnalysis(Positive["reviews"], 1, 20) Bigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Bigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Bigram of reviews with Positive Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  #Neutral words = GramAnalysis(Neutral["reviews"], 1, 20) Bigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Bigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Bigram of reviews with Neutral Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  #Negative words = GramAnalysis(Negative["reviews"], 1, 20) Bigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Bigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Bigram of reviews with Negative Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show() </pre>

	<pre># trigram - Positive, Negative, Neutral  #Positive words = GramAnalysis(Positive["reviews"], 1, 20) Trigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Trigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Trigram of reviews with Positive Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  #Neutral words = GramAnalysis(Neutral["reviews"], 1, 20) Trigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Trigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Trigram of reviews with Neutral Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()  #Negative words = GramAnalysis(Negative["reviews"], 1, 20) Trigram = pd.DataFrame(words, columns = ["words", "counts"])  #Visualization Trigram.groupby("words").sum()["counts"].sort_values().plot(kind = "barh", color = "green", figsize = (10,5)) plt.title("Trigram of reviews with Negative Sentiments") plt.xlabel("Total Counts", color = "magenta", fontsize = 10, labelpad = 15) plt.xticks(rotation = 0) plt.ylabel("Top Words", color = "cyan", fontsize = 10, labelpad = 15) plt.show()</pre>
--	---

Model 4 – Training the Models

```
#Splitting dataset
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_final, y_final, test_size=0.25, random_state=42)

# Model Selection & Evaluation

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

LogReg = LogisticRegression()
RForest = RandomForestClassifier()

LogReg.fit(X_train, y_train)

LogReg.predict(X_test)

RForest.fit(X_train, y_train)

RForest.predict(X_test)

y_pred_LogReg = LogReg.predict(X_test)
y_pred_RForest = RForest.predict(X_test)

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

accuracy_LogReg = accuracy_score(y_test, y_pred_LogReg)
accuracy_RForest = accuracy_score(y_test, y_pred_RForest)
```

Model 5 – Evaluation on Test Set

```
print("Accuracy of Logistic Regression:", accuracy_LogReg)
print("Accuracy of Random Forest:", accuracy_RForest)

Accuracy of Logistic Regression: 0.9524161371361816
Accuracy of Random Forest: 0.9757647406531698

cm = confusion_matrix(y_test, y_pred_RForest)
cm

array([[2201, 0, 24],
       [ 0, 2226, 51],
       [ 7, 82, 2176]])

# Classification Scores
print(classification_report(y_test, y_pred_RForest))

precision recall f1-score support

0 1.00 0.99 0.99 2225
1 0.96 0.98 0.97 2277
2 0.97 0.96 0.96 2265

accuracy 0.98 6767
macro avg 0.98 0.98 0.98 6767
weighted avg 0.98 0.98 0.98 6767
```

## Model 6 – Individual Review Predictions

```
# --- Step 1: Create a function for text cleaning and preprocessing ---
def preprocess_new_text(text):
    # Clean text (use the same functions you defined earlier)
    text = Text_cleaning(text)
    text = Text_preprocessing(text)
    return text

# --- Step 2: Function to predict sentiment for new text ---
def predict_sentiment(review, model, vectorizer, encoder):
    # Preprocess input review
    processed_review = preprocess_new_text(review)

    # Transform using the trained TF-IDF vectorizer
    vectorized_review = vectorizer.transform([processed_review]).toarray()

    # Predict sentiment
    pred = model.predict(vectorized_review)

    # Decode label back to sentiment
    sentiment_label = encoder.inverse_transform(pred)[0]
    return sentiment_label

# --- Step 3: Example usage ---
sample_review_1 = "This instrument has amazing sound quality and I love it!"
sample_review_2 = "The product is terrible, very disappointed."
sample_review_3 = "It is okay, not the best but not the worst either."

print("Review:", sample_review_1, "-> Sentiment:", predict_sentiment(sample_review_1, Rforest, TF_IDf, encoder))
print("Review:", sample_review_2, "-> Sentiment:", predict_sentiment(sample_review_2, Rforest, TF_IDf, encoder))
print("Review:", sample_review_3, "-> Sentiment:", predict_sentiment(sample_review_3, Rforest, TF_IDf, encoder))
```

Review: This instrument has amazing sound quality and I love it! -> Sentiment: Positive  
Review: The product is terrible, very disappointed. -> Sentiment: Positive  
Review: It is okay, not the best but not the worst either. -> Sentiment: Positive